

# pig, the PAG interface generator

Gergő Bárány

February 20, 2005

## 1 Introduction

**pig** is the PAG interface generator; it is a tool for automatical generation of large parts of the glue code between a compiler front end and PAG, the program analyzer generator. More precisely, **pig** is designed to generate the AST access functions required by PAG (these are the `is_op` functions for type tests, the `get` functions for child access, and several functions for syntactic list traversal). Due to its flexible specification language, **pig** can also be used to generate such access functions for other applications.

**pig** is specialized for AST access since, given a well-designed AST, most such accesses share a common structure and only differ in details such as the names of the fields to be accessed or types to be compared. This structure and the fact that many such functions must be present suggest an automated approach to code creation. In contrast, **pig** cannot be utilized for generating the CFG access functions, as each of these functions is unique and looks different from the others.

The following sections describe in detail how **pig** operates and how you can operate it.

## 2 Input and output

The input to **pig** consists of two files, one specifying the abstract syntax and one containing macro rules describing the functions to be created. It generates a C header and a code file. This is sketched in Figure 1 below.

The first input file is called the *syn file*. It should be identical to the `syn` file required by PAG. Thus, it contains a tree grammar describing the structure of the abstract syntax trees for which code is to be generated. It may also contain equations declaring alias types; these are ignored by **pig**. The syntax of the `syn` file is not described further, you can find a description in the PAG manual.

The other input file is the *spec file*. It contains rules that describe which functions to create and how to create them. This is the file against which all identifiers in the `syn` file are matched; its syntax and the semantics of the matching and code generation are described in the following section.

The names of the `syn` and `spec` files may be chosen freely, they are passed to **pig** on the command line.

The output files are called `syntree.h` and `syntree.c` as required by PAG. The header file contains prototypes for all functions generated by **pig** or alternatively the macro




Figure 1: Inputs and outputs of `pig`

definitions for the functions that are to be implemented as macros. The code file contains the definitions of the real functions specified by the user. Additionally, `pig` can add user-defined pre- and postfixes to both files. This is useful for `#include` statements, `typedefs`, auxiliary functions and the like.

### 3 Spec file syntax

The syntax of the spec file is given below in an EBNF-like notation, using `[]` for options, `|` for alternatives, and `*` for zero or more repetitions. Nonterminals are typeset in *italics*, terminal symbols in **typewriter** type. Additional tokens are *Code* for C code fragments, anything (except `%}`) enclosed in `%{` and `%}`; *Id* for an identifier, either a single `_` or a letter followed by any number of letters, digits, or `_`; *Str* for a string, an identifier enclosed between `"`; and *Num* for a number, a non-empty sequence of digits.

Whitespace and comments (the usual types with `//` and `/* ... */`) may be inserted anywhere between tokens.

```

spec      ::= Code Code section* Code Code
section   ::= %% keywords head body*
keywords ::= [per_constructor] [list] [macro]
head      ::= Id : head (Id, Id, Id, Id, Id, Id, Id) Code
body      ::= Id : body (Id, idstr, idstr, idnum, idstr, idnum, idstr) Code
idstr     ::= Id | Str
idnum     ::= Id | Num

```

Each section is assigned a name, the identifier in front of the `head` keyword. Each `body` rule in that section must be assigned the same name.

Inside the argument lists of head and body rules, the identifiers play the role of variable names, while strings and numbers are constants. These argument lists are used for

matching identifiers from the syn file to rules in the spec file, as described in the following section.

## 4 Rule matching and the spec file

The syn file consists of productions of the following form:

```

type:  constructor(field: field type, ..., field: field type)
      :
      |  constructor(field: field type, ..., field: field type)
      ;

```

**pig**'s matcher traverses all such productions, and for each field occurring in each production's alternatives, it creates a tuple

*(type, constructor, constructor index, field, field index, field type).*

The constructor index is  $j$  for the  $j$ -th constructor for this type, the field index is  $i$  for the  $i$ -th field for the current constructor. Both are counted from 0.

The spec file is divided into sections, each of which describes one function or macro to be implemented. A section consists of several rules. The first rule in the section, called the *head* rule, describes the function's prototype (or macro head), the other rules are *body* rules describing possible implementations of the function.

Each body rule is associated with an argument list whose members (except the first one, which is ignored in matching) correspond to the members of the tuples created from the syn file as described above; these can be variables or constants of the appropriate type (C identifiers for the type, constructor, field, and field type; nonnegative integers for the indices).

The tuples from the syn files are matched against these rules: A variable matches anything, a constant matches an identical constant. In particular, an argument list containing only variables will match every tuple. **pig** chooses the first body rule matching the constants from the syn file. The head and this body are then pasted together and all variables instantiated with the constants; the result is a function or macro definition that is emitted into the output files.

Variable substitution is performed by the C preprocessor whose **##** operator can also be used to paste tokens together.

The semantics of the three keywords are the following: The **per\_constructor** keyword means that the function defined in this section should not be instantiated for every field but only once for each constructor. This is necessary for functions that do not depend on the field, for instance the **is\_op** functions required by PAG. If **per\_constructor** were not used, **pig** would emit several identical function definitions, causing problems at compile time.

**list** means that the section should only be instantiated for list fields, i.e. ones marked with **\*** in the syn file. This is used for list traversal functions such as **empty**, **hd** and **tl**. Implementing such functions for non-list types would not make sense.

Finally, **macro** instructs **pig** to create the definitions for this section not as functions but as C preprocessor macros.

## 5 An example

These theoretical explanations are best illustrated by an example. Consider the following fictitious fragment of a syn file:

```
Statement: Assign(left: Var, right: Expr)
          | Loop(cond: Expr)
          ;
Expr      : Binary(left: Expr, right: Expr)
          | Unary(child: Expr)
          ;
```

Let us assume that the fields of a **Statement** node are to be accessed by indexing an array of pointers, while the children of **Exprs** are accessed by name. Type information for the **is\_op** functions is stored in an enumeration. Additionally, the **get** functions should be implemented as macros to avoid the overhead of function calls.

The spec file might look like this:

```
%{ // prefix of syntree.c
#include "syntree.h"
%}

%% // section get
macro get:head(NODE, TYPE, CONSTR, _, FIELD, FIELDI, FTYPE)
%{
TYPE##_##_CONSTR##_get_##_FIELD(NODE)
%}

get:body(NODE, "Statement", _, _, _, FIELDI, _)
%{
    (NODE->children[FIELDI])
%}

get:body(NODE, "Expr", _, _, FIELD, _, _)
%{
    (NODE->FIELD)
%}

%% // section is_op
per_constructor is:head(NODE, TYPE, CONSTR, _, _, _, _)
%{
int is_op_##_TYPE##_##_CONSTR(TYPE NODE)
%}

is:body(NODE, _, CONSTR, _, _, _, _)
%{
    return (NODE->type == T_##_CONSTR);
%}
```

```
%{%} // empty postfixes for syntree.h and syntree.c
%{%}
```

Given these inputs, `pig` produces two files, `syntree.h` and `syntree.c`. The contents of `syntree.h` are the following:

```
// prefix of syntree.h
#include "type_definitions.h"

#define Statement_Assign_get_left(NODE) (NODE->children[0])
int is_op_Statement_Assign(Statement NODE);
#define Statement_Assign_get_right(NODE) (NODE->children[1])
#define Statement_Loop_get_cond(NODE) (NODE->children[0])
int is_op_Statement_Loop(Statement NODE);
#define Expr_Binary_get_left(NODE) (NODE->left)
int is_op_Expr_Binary(Expr NODE);
#define Expr_Binary_get_right(NODE) (NODE->right)
#define Expr_Unary_get_child(NODE) (NODE->child)
int is_op_Expr_Unary(Expr NODE);
```

Note that the `#define` statements are intermixed with function prototypes. The reason for this is that `pig` emits definitions grouped by the fields defined in the `syn` file rather than by the sections of the `spec` file. Observe also how the `##` operator is used to create single tokens such as `is_op_Expr_Binary` out of multiple parts and macros.

The `syntree.c` file looks like this (slightly reformatted to fit the page):

```
// prefix of syntree.c
#include "syntree.h"

int is_op_Statement_Assign(Statement NODE)
{ return (NODE->type == T_Assign);}
int is_op_Statement_Loop(Statement NODE)
{ return (NODE->type == T_Loop);}
int is_op_Expr_Binary(Expr NODE){ return (NODE->type == T_Binary);}
int is_op_Expr_Unary(Expr NODE){ return (NODE->type == T_Unary);}
```

The function definitions are completely expanded; this can be requested with a flag passed to `pig`. The default behavior is to only emit a macro definition for each function and macro invocations for each concrete function to be defined. This is a little more compact and might be more readable under certain circumstances.

In this small example, the specification for `pig` is longer than the generated code. In real-world applications, however, the `syn` file is considerably larger than in this case, causing much more code to be emitted. Additionally, having a simple specification all function definitions adhere to will often ease debugging, maintenance and the addition of new features.